


# Introduction to IP Routing


# Routing

- How do packets get from A to B in the Internet?



# Connectionless Forwarding

- Each router (switch) makes a LOCAL decision to forward the packet towards B



# Connectionless Forwarding

- This is termed *destination-based connectionless forwarding*
- How does each router know the *correct* local forwarding decision for any possible destination address?
  - Through knowledge of the *topology state* of the network
  - This knowledge is maintained by a *routing protocol*

# Routing Protocols

- Distribute the knowledge of the current *topology state* of the network to all routers
- This knowledge is used by each router to generate a *forwarding table*, which contains the local switching decision for each known destination address

# Forwarding Table


---

```
nancy@sluggo.lab> show route forwarding-table
```

```
Internet:
```

Destination	Type	RtRef	Nexthop	Type	Index	NhRef	Netif
10.100.71.0/24	user	0	10.100.67.254	ucst	18	74212	GigE0.0
10.100.71.224/27	user	2	10.100.67.254	ucst	18	74212	GigE0.0
10.250.1.36/30	intf	0	ff.3.0.21	ucst	27	1	so-2/0/0.0
10.250.1.37/32	intf	0	10.250.1.37	locl	26	1	
10.250.1.103/32	dest	0	10.250.1.103	bcst	37	1	ge-7/2/0.0

--- (more) ---




# Routing Protocols

- correct operation of the routing state of a network is essential for the management of a *quality* network service
  - accuracy of the routing information
  - dynamic adjustment of the routing information
  - matching aggregate traffic flow to network capacity

# ISP Routing Tasks

- customers
- internal
- peer / upstream






# Interior Routing


- discovers the *topology* of a network through the operation of a *distributed routing protocol*

# Path Selection



Minimum cost from A to B is **39** units

# Dynamic Path Adjustment



If R5 – R7 breaks, minimum cost path from A to B is Now **46** units

# Interior Routing Protocols

- describe the current network topology
- Routing protocols distribute **how** to reach **address prefix** groups
- Routing protocols function through either
  - *distributed computing model (distance vector)*
  - *parallel computing model (link state)*

# Routing Protocols

- Distance Vector Routing Protocols
  - Each node sends its routing table (dest, distance) to all neighbors every 30 seconds
  - Lower distances are updated with the neighbor as next hop
  - cannot scale
  - cannot resolve routing loops quickly
  - RIP is the main offender

# Routing Protocols

- Link State Routing Protocols
  - Each link, the connected nodes and the metric is flooded to all routers
  - Each link up/down status change is incrementally flooded
  - Each router re-computes the **routing table** in parallel using the common **link state database**.
  - OSPF is the main protocol in use today

# Suggestions

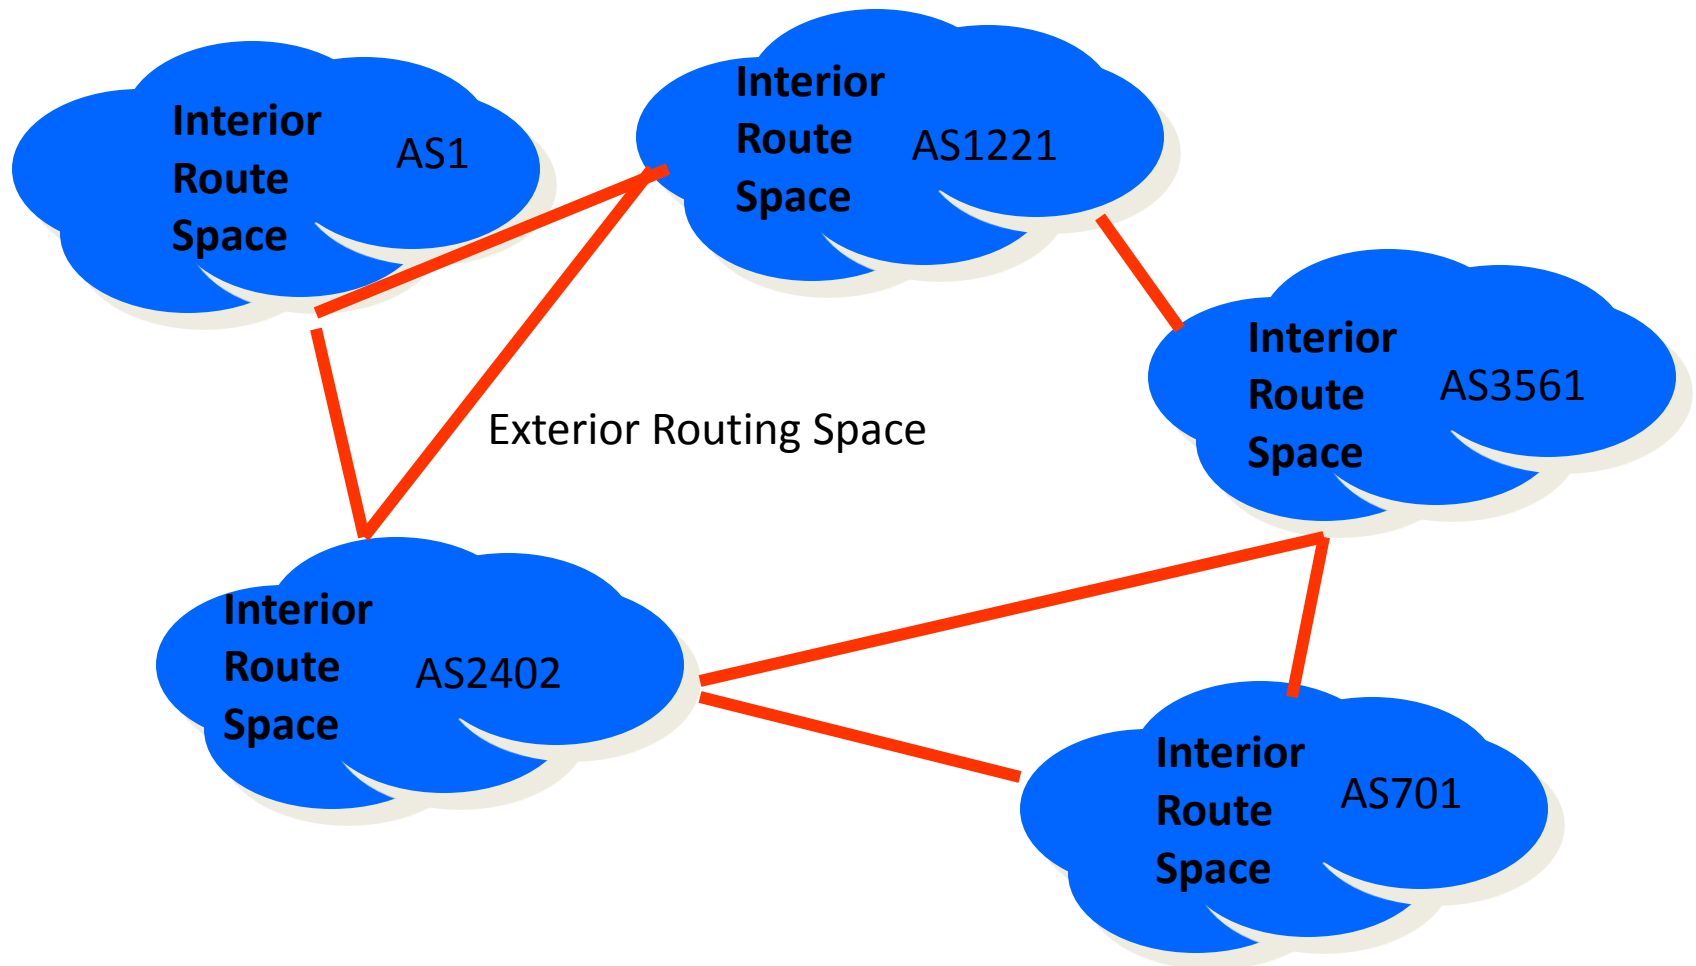
- Just engineering a physical link does not ensure that traffic will flow
  - some system somewhere must provide routing information about how to reach the newly connected network
- Installing backup circuits is easy, making the routing work may not be

# Suggestions

- need a clear understanding of how the client networks want their traffic to flow before you can start making routing configuration changes



# Interior and Exterior Routing Protocols




# Exterior Routing Protocols

- You tell me all the **address prefixes** you can reach, but don't tell me the path you use to get there
  - I'll tell you the same
- If anything changes, please let me know
- If you tell me an address I'll send you traffic destined to that address.
  - If I tell you an address I will accept traffic destined to that address

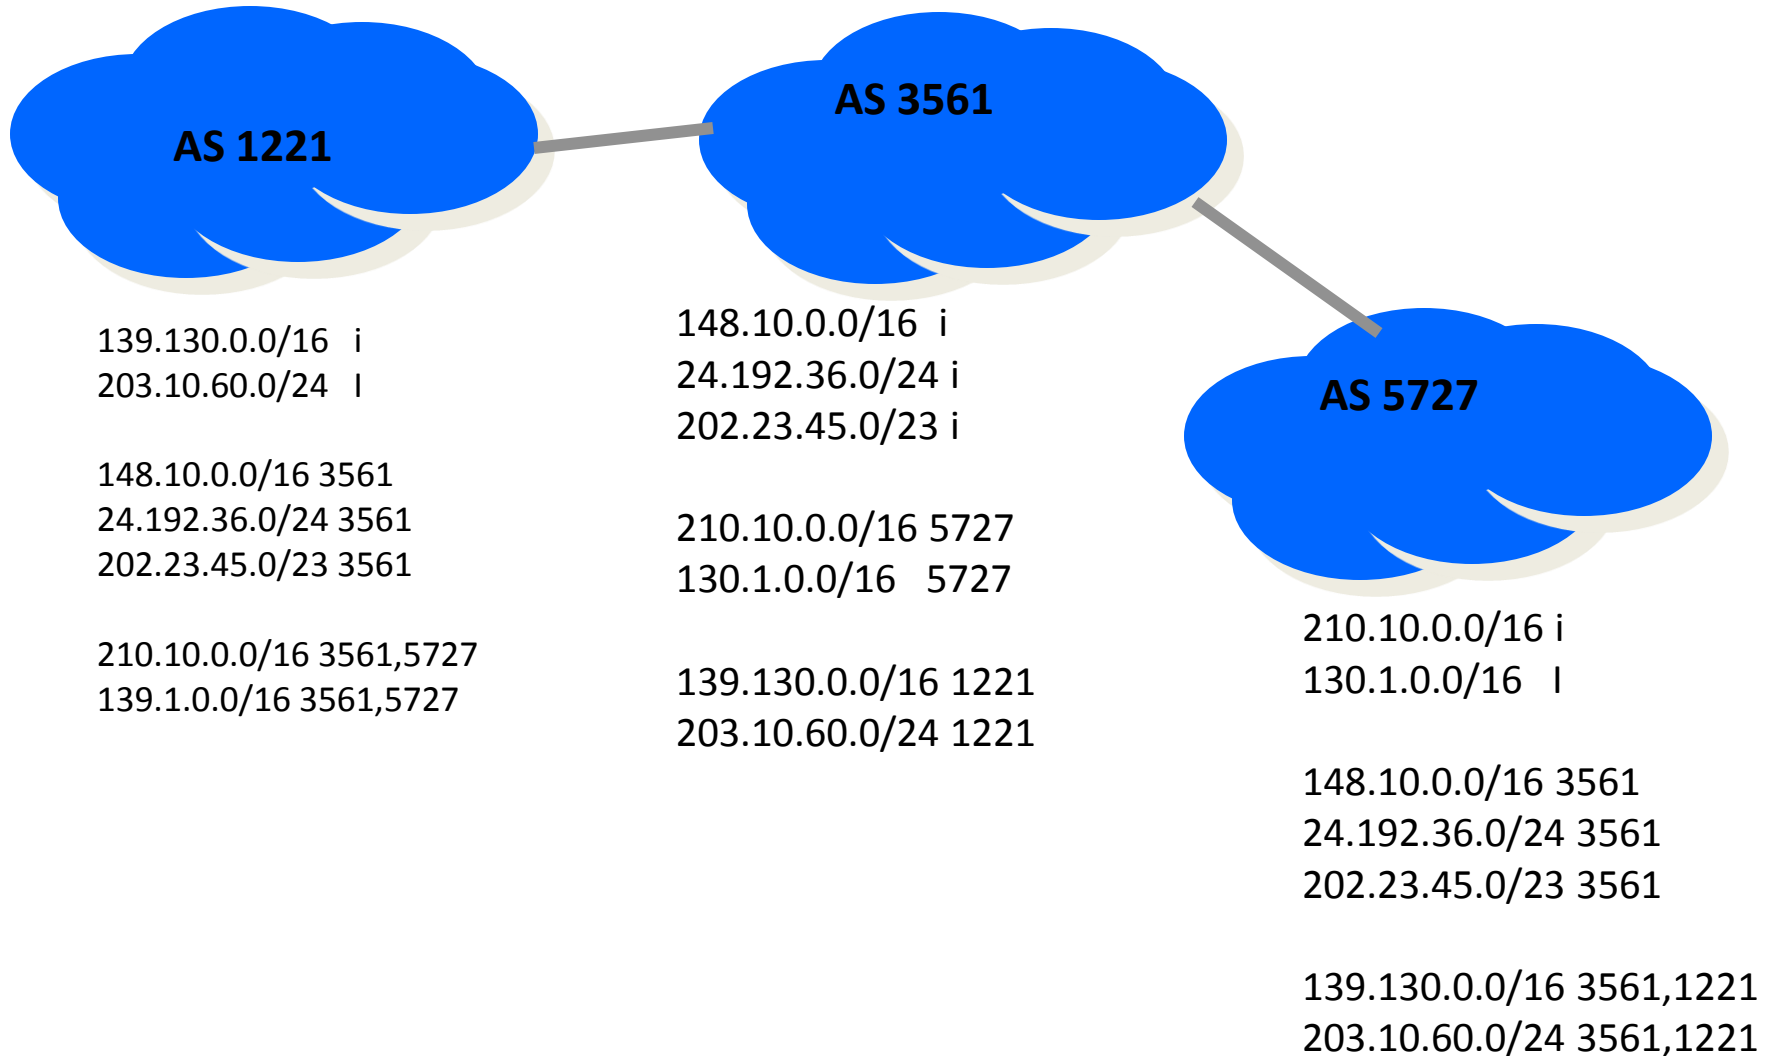
# Exterior Routing Protocols

- Border Gateway Protocol version 4 (BGP4)
- Each interior route collection is described by an **Autonomous System (AS)** number
- Internal topology is hidden
- Routes are announced with associated AS value
  - 139.130.0.0/16 + AS 1221

# BGP example




# BGP Example of TRANSIT



# Exterior Routing Protocols


- Internal transit paths use I-BGP



Q: How does router A tell router B about AS1221 addresses?

A: Router A sets an INTERIOR BGP session with router B

# Exterior Routing Protocols



# Exterior Routing Protocols

- Normally chose minimal AS path length

—————→ 203.10.60.0/24 701,3561,1221  
                  203.10.60.0/24 5727,1221


Selected path is via peer session to AS 5727 as this  
Is 1 AS shorter than the other path



# Exterior POLICY

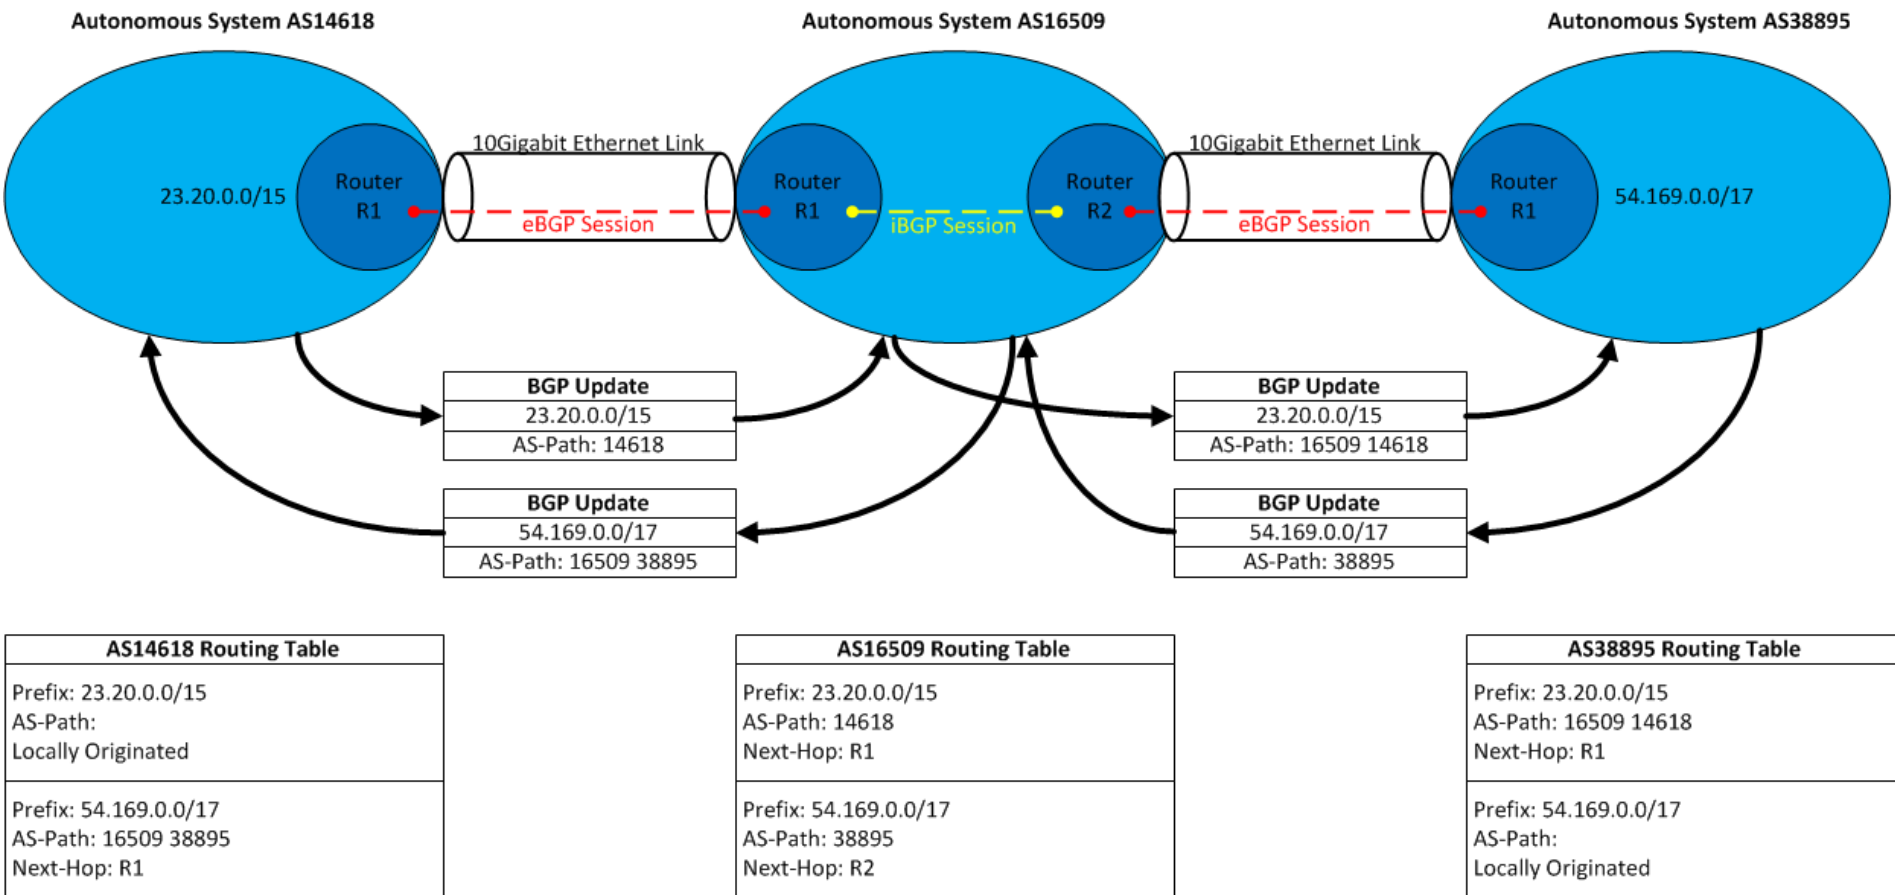
- How can I share the traffic load between 2 or more exterior providers?
- How can I create a backup link to support my main exterior link?
- You can bias minimal path selection by **AS path filter lists** or **community attributes** or **local preferences**

# Exterior Routing Protocols plus Policy



# Exterior Routing Protocols plus Policy

- policy settings control
  - what you advertise to your immediate peers
  - What you accept from your immediate peers
  - What transits you will accept (send traffic)
- you **cannot** control
  - transit path of received traffic
  - symmetry of transit policy



# Interior Gateway Protocols

## Routing Information Protocol (RIP)

- IP only
- distance vector protocol
- slow convergence
- does not carry mask information
- reasonably simple design & configuration
- does not scale (maximum 15 hops)
- poor metrics (hop-count)

# Interior Gateway Protocols

## The IGRP metric

»always get optimal routing metric vector, not single value

- I. bandwidth
- II. delay
- III. hops
- IV. reliability
- V. loading

# Interior Gateway Protocols

## Open Shortest Path First (OSPF)

- I. IP only
- II. link state protocol
- III. fast convergence
- IV. design and architecture very complex
- V. configuration can be simple

# Interior Gateway Protocols

## Which to use?

- a. Your interior network is actually VERY simple.
- b. Your IGP should only carry your routes and your direct customers'



# Interior Gateway Protocols

Problems with "classic" protocols

- ❑ slow convergence
- ❑ count to infinity
- ❑ no mask information



# Interior Gateway Protocols

## *Slow convergence*

advertisement period

- entire routing table dumped every n seconds

timeout period

- usually 3 times advertisement period

RIP values are normally 30 and 90 seconds!

## Count to infinity: hold-down

**Link Between A & B is Broken**



	A	B	C	D
A	0, -	1, A	2, B	3, C
B	1, B	0, -	2, C	3, D
C	2, B	1, C	0, -	1, D
D	3, B	2, C	1, D	0, -

# Count to infinity problem

- ☐ Imagine a network with a graph as above.
- ☐ As you see in this graph, there is only one link between A and the other parts of the network.
- ☐ Now imagine that the link between A and B is cut.
- ☐ At this time, B corrects its table.
- ☐ After a specific amount of time, routers exchange their tables, and so B receives C's routing table.
- ☐ Since C doesn't know what has happened to the link between A and B, it says that it has a link to A with the weight of 2 (1 for C to B, and 1 for B to A -- it doesn't know B has no link to A).
- ☐ B receives this table and thinks there is a separate link between C and A, so it corrects its table and changes infinity to 3 (1 for B to C, and 2 for C to A, as C said).
- ☐ Once again, routers exchange their tables.
- ☐ When C receives B's routing table, it sees that B has changed the weight of its link to A from 1 to 3, so C updates its table and changes the weight of the link to A to 4 (1 for C to B, and 3 for B to A, as B said).
- ☐ This process loops until all nodes find out that the weight of link to A is infinity.