

CO324: TCP clients and servers

Ziyan Maraikar

February 12, 2015

Lecture Outline

1 TCP Clients

2 TCP Servers

3 Framing

TCP

TCP is a *stream protocol* in which packet boundaries are invisible to the application. Data is received and transmitted as a sequence of bytes.

It provides

- ① Reliability
- ② Ordering
- ③ Flow control

TCP clients

The Java `Socket` class represents a TCP socket.

```
Socket socket = new Socket();
```

```
InputStream sin = socket.getInputStream();
```

```
OutputStream sout = socket.getOutputStream();
```

Binary data is read and written to the socket via the associated I/O streams.

Sending and receiving text

```
try (Socket socket = new Socket(address, TCPServer.PORT);
    Scanner sin = new Scanner(socket.getInputStream() );
    PrintStream sout = new PrintStream(
        socket.getOutputStream() ) ) {

    sout.println("client says hello to server "+address);
    System.out.println(sin.nextLine());
}
```

Note that application messages must be properly *framed* e.g. using a delimiter like `\n`.

What do the `Scanner` and `PrintStream` classes do?

Lecture Outline

- 1 TCP Clients
- 2 TCP Servers**
- 3 Framing

TCP Servers

Java uses a separate `ServerSocket` class to bind to a port and accept connections from clients.

```
ServerSocket ss = new ServerSocket(PORT);  
Socket socket = ss.accept();
```

`accept` returns a new socket connected to the client.

Server example

```
try( ServerSocket ss = new ServerSocket(PORT)) {  
    while (true)  
        try (Socket socket = ss.accept();  
            Scanner sin = new Scanner(socket.getInputStream() );  
            PrintStream sout = new PrintStream(  
                socket.getOutputStream()) ) {  
  
                sout.println("server says world");  
                System.out.println(sin.nextLine());  
            }  
        }  
}
```

What happens if multiple clients try to connect at once?

Lecture Outline

- 1 TCP Clients
- 2 TCP Servers
- 3 Framing**

Messages on streams

Suppose a client sends two consecutive messages, and the server does a `read`. Will it get

- ① both messages at once?
- ② the first message only?
- ③ part of the first message?

It depends on network conditions and the TCP/IP stacks!

Framing

TCP can only send to and receive from a byte stream, but application protocols are built with discrete messages.

We must define a method of *framing* application messages, so that message boundaries are unambiguous.

Method used depends on the kind of data

- ★ Text
- ★ Binary

Text protocols

Most application protocols on the Internet are textual.

- ★ Human readable — so easy to debug.
- ★ Historically, the most applications were textual e.g., Telnet, Email

Disadvantages:

- ★ Vulnerable to security attacks like buffer overflows.
- ★ Bandwidth and CPU inefficient.

Example: SMTP

```
S: 220 smtp.server.com Simple Mail Transfer Servi
C: HELO client.example.com
S: 250 Hello client.example.com
C: MAIL FROM:<jane@yahoo.com>
S: 250 OK
C: RCPT TO:<john@gmail.com>
S: 250 OK
C: DATA
S: 354 Send message content; end with <CRLF>.<CRL
C: <The message data (body text, subject, e-mail
C: .
S: 250 OK, message accepted for delivery: queued
C: QUIT
S: 221 Bye
```

How are non-text mail attachments handled?

Delimiters

We can *delimit* text protocol messages using a special character. The usual delimiter used in Internet protocols are the line termination characters **CR**, **LF** or **CRLF**.

```
Socket socket = new Socket(address, PORT);

BufferedReader sin = new BufferedReader (
    new InputStreamReader(socket.getInputStream() ));

BufferedWriter sout = new BufferedWriter (
    new OutputStreamWriter(socket.getOutputStream() ));

sout.write("hello world\n");
sin.readLine();
```

BufferedReader.readLine splits the apart newline delimited messages in a stream.

Binary protocols

Binary protocols support transmission of arbitrary data. Usually contains a fixed-format *header* that describes the *payload*.

- ★ Suited to describing structured data.
- ★ Easier to *parse* — metadata received before payload.
- ★ Efficient use of bandwidth.

Example: Basic encoding rules for ASN.1, an OSI standard used in protocols such as LDAP.

Type	Length	Value	End-of-content
------	--------	-------	----------------